

MSc DevOps, Software Evolution and Software Maintenance — Final Report

Group i — *I-Terroni-DevOps* Members: Michael Fantinato, Vincenzo Sabino, Gabriele Matteoli, Rachele Russo, Benedek Szabo **Repository:** <https://github.com/stegish/I-Terroni-DevOps> **Date:** May 2026

Linked Artifacts

Artifact	Link / Location
Main repository	https://github.com/stegish/I-Terroni-DevOps
Issue tracker	https://github.com/stegish/I-Terroni-DevOps/issues
Production application	http://164.92.231.30/public (DigitalOcean)
CI/CD pipelines	.github/workflows/continuous-deployment.yml , .github/workflows/code-quality.yml , .github/workflows/build-report.yml
Container images	michaelfant/minitwitimage:latest , michaelfant/flagtoolimage:latest on Docker Hub
Infrastructure as Code	infrastructure/docs/infrastructure-as-code.md
Security report	SECURITY.md
Grafana dashboards	monitoring/grafana/dashboards/
Code-quality dashboards	SonarCloud + Codacy projects

1. System's Perspective

1.1 Design and Architecture

ITU-MiniTwit is a Twitter-like micro-blogging service that we rewrote from the given Flask + raw-SQL code base into a **Pyramid** application with **SQLAlchemy ORM** data layer. The web framework choice is documented in [README.md](#): We chose Pyramid instead of Bottle or Flask because it gave us a clearer structure for the application. Bottle would have required more manual setup for things like templating, while Flask often depends on global application state. Pyramid uses an explicit request object, which made it easier for us to attach things like the database session to `request.db`. This also made the code easier to test, since individual request handlers could be tested without needing to rely on a full application context.

In production the system is a **3-node Docker Swarm** running on DigitalOcean droplets:

- **1 manager (s-2vcpu-2gb)**: hosts the observability stack (Prometheus, Grafana, Loki) and the nginx ingress. It is deliberately kept off the application path so the app and the telemetry never compete for RAM.
- **2 workers (s-1vcpu-1gb each)**: run **3 minitwit replicas** (Pyramid + 3 gunicorn workers each) plus the **flagtool** admin container.
- Per-node agents (**node-exporter**, **cadvisor**, **promtail**): run in Swarm **mode: global** one task per droplet.

The database runs as a separate DigitalOcean Managed MySQL 8 instance outside the Docker Swarm. The application connects to it through the `DATABASE_URL` environment variable, using TLS for the connection. On the manager node, nginx handles HTTPS traffic and forwards requests into the Swarm network to the MiniTwit service through Swarm DNS (`tasks.minitwit:5000`). We also use this DNS-based service discovery for monitoring, so Prometheus can scrape the actual running replicas instead of only reaching one container through a single load-balanced address.

Deployment safety is enforced through a mandatory healthcheck: each replica runs a Python `urllib` probe every 10 seconds, with a 20-second grace period before Swarm considers the container ready. This prevents the **start-first** rolling strategy from terminating an old replica before the new one is fully initialized and accepting connections. Additionally, services declare varying restart policies: the main application uses **condition: any** (restart on all exits), while observability and utility services use **condition: on-failure** with exponential backoff (delay, max_attempts, window), allowing graceful shutdown and containment of transient failures.

1.2 Dependencies

Layer	Tooling
Language / runtime	Python 3.12-slim
Web framework	Pyramid + <code>pyramid_jinja2</code> (templating)
WSGI server	<code>gunicorn</code> (3 workers × 2 threads)
ORM / DB driver	SQLAlchemy + PyMySQL
Database	MySQL 8 (DigitalOcean Managed external instance, TLS connection via <code>DATABASE_URL</code> environment variable); SQLite in CI
Frontend	server-rendered Jinja2 templates + static CSS

Layer	Tooling
Observability	Prometheus 2.55, Grafana 11.4, Loki 2.9, Promtail 3.0, node-exporter 1.8, cAdvisor 0.49, mysqld-exporter 0.14.0, prometheus-client ; per-node agents run <code>mode: global</code> with resource caps to prevent interference with application workloads
Reverse proxy / TLS	nginx 1.27-alpine + Let's Encrypt (certbot)
Container orchestration	Docker Engine + Docker Swarm (compose-spec)
Infrastructure as Code	Terraform (DigitalOcean provider); Vagrantfile kept for single-node local experiments only
CI/CD	GitHub Actions; Docker Hub registry
Static analysis	ruff , codespell , mypy , hadolint , shellcheck
Security scanning	Semgrep (SAST), Trivy (image CVEs), SonarCloud, Codacy
Browser E2E tests	Selenium (standalone-chrome)




Figure 1: Dependency graph: logical view

Figure 1.2.1: Dependency graph: logical view

The full dependency manifest is in `requirements.txt` and the lint/format/type-

check configuration in `pyproject.toml`. The CI image is built from `Dockerfile-minitwit-tests` and never reaches production, keeping test dependencies and debug code out of `michaelfant/minitwitimage:latest`.

1.3 Current state, static analysis & quality assessment

Quality is continuously measured by two third-party services hooked to every push and PR (workflow: `.github/workflows/code-quality.yml`):

- **SonarCloud** reports maintainability, reliability and security ratings, duplications, cyclomatic complexity and the SQALE technical-debt index.
- **Codacy** provides an aggregated grade combining `ruff`, `pylint`, `bandit`, `hadolint` and `shellcheck`.

At hand-in time both projects sit at a passing quality gate. Local linting passes cleanly (`make check`), `ruff format --check` is green, and the latest Trivy scan of `michaelfant/minitwitimage:latest` reports no HIGH/CRITICAL findings with `ignore-unfixed: true` (the gate that blocks deploy). `mypy` runs non-blocking and surfaces residual type gaps in legacy modules.

2. Process' Perspective

2.1 CI/CD pipeline

The pipeline operates in `.github/workflows/continuous-deployment.yml` and is structured as four sequential jobs, each a hard gate for the next:

`static-analysis` → `tests` → `security-scan` → `build-and-deploy`

1. **static-analysis:** `ruff` (lint + format), `codespell`, `mypy` (non-blocking), `hadolint` on the three Dockerfiles, `shellcheck` on `control.sh/deploy.sh`, and **Semgrep SAST** with the `p/security-audit`, `p/owasp-top-ten`, `p/python` and `p/dockerfile` rule packs, failing the build on findings of severity `ERROR`.
2. **tests:** builds the production image locally, spins up MySQL 8 + Selenium Chrome on a Docker network, runs the schema init (mirroring `deploy.sh`), and then executes the three test suites: integration (`minitwit_tests_refactor.py`), simulator API (`minitwit_sim_api_test.py`) and Selenium (`test_itu_minitwit_ui.py`).
3. **security-scan:** `trivy` scans the built image for OS-package and Python-dependency CVEs and fails on HIGH/CRITICAL. Results are also uploaded as SARIF to the GitHub Security tab.
4. **build-and-deploy:** only runs on `push` to `main` (not on PR). Pushes the two production images to Docker Hub and SSH into the manager droplet to run `deploy.sh`, which uses `docker stack deploy` with a rolling update.




Figure 2: CI/CD pipeline: sequential gates

Figure 2.1.1: CI/CD pipeline: sequential gates

Infrastructure is provisioned with **Terraform** (`infrastructure/main.tf`). Remote-exec provisioners automate `docker swarm init` on the manager and `docker swarm join` on workers using a join token retrieved via SSH.

The deploy itself does three important things beyond `docker stack deploy`: it computes a `sha256` hash of the `nginx` config and `promtail` config and injects them as labels / Swarm config names so that a config-only change reliably triggers a rolling restart, and it runs `init_db()` exactly once in a container.

2.2 Monitoring

We collect metrics with Prometheus and visualise them in six domain oriented Grafana dashboards (01-business, 02-api-http, 03-system-health, 04-infrastructure, 05-database, 06-logs), provisioned automatically from `monitoring/grafana/dashboards/`.

What we monitor:

- **Business:** total users, total messages, total follow relations, average followers, registration/message rate, derived from in-app gauges in `metrics.py`.
- **API / HTTP:** request count by `method × route × status`, latency histogram, 4xx/5xx rates, error budget burn.
- **System health:** droplet CPU, memory, disk and network from `node-exporter`; per container CPU/RAM/IO from `cAdvisor`.
- **Infrastructure:** per droplet resource health (CPU, memory, disk, network utilization, load average) with per node breakdowns and thresholds.
- **Database:** InnoDB metrics, process list, query throughput, table-lock waits from `mysqld-exporter` against the DO managed MySQL.
- **Logs:** Loki query panels showing 5xx bursts and recent error lines.

Prometheus discovers replicas via Swarm overlay DNS (`tasks.minitwit`, `tasks.node-exporter`, `tasks.cadvisor`) so every replica/agent is scraped. TSDB retention is capped at **7 days or 512 MB**, whichever comes first, so the manager droplet never runs out of disk.

2.3 Logging

Promtail runs and tags each line with the container name and delivers to Loki on the manager (see `logging/promtail-config.yml`). The application uses `json-log-formatter` so stdout lines are structured JSON, which makes Loki labels actually useful. Retention is 7 days, enforced by the Loki compactor. Grafana’s “06-logs” dashboard exposes a small LogQL toolbox for quick incident triage.

2.4 Security hardening

A full risk assessment + mitigation plan can be found in `SECURITY.md`. Headline items, all implemented in this branch:

- **Firewall defense in depth.** Docker rewrites `iptables` and bypasses `ufw`, so we layered two controls: (a) `ufw` on the host (allows only 22/80/443) and (b) DigitalOcean cloud firewall declared in `infrastructure/firewall.tf`. Internal ports (Prometheus 9090, Grafana 3000, Loki 3100) are reached over SSH tunnels only.
- **TLS** via `nginx` + Let’s Encrypt; all external traffic is HTTPS-only, but internal traffic between `nginx` and the app is plain HTTP (HTTP redirects to HTTPS); renewal is in cron.
- **Non-root containers.** All three Dockerfiles add a dedicated `appuser` and `USER appuser` before `CMD`.
- **Secrets out of source.** Removed hard-coded simulator credentials and the default Pyramid `SECRET_KEY`; both are now required env vars that crash the app on startup if missing.
- **Image base bump** `python:3.9-slim` → `python:3.12-slim`.
- **Shift-left in CI:** Semgrep (SAST) and Trivy (image CVE) gate the deploy.
- **Third-party Actions pinned by commit SHA** (e.g. `SonarSource/sonarcloud-github-action@ffc3`) so a hijacked tag can’t exfiltrate workflow secrets.

2.5 Availability and scaling

Availability comes from three layers:

1. **Replica horizontality.** `minitwit` runs as 3 Swarm replicas, so a worker droplet can die and Swarm keeps at least one replica serving on the survivor. The `update_config: { order: start-first }` guarantees we never go below 3 replicas during a rolling deploy.
2. **Resource limits.** Every service declares both `reservations` and `limits` (e.g. `minitwit` 128 MB reserved / 256 MB cap). A misbehaving container cannot run the droplet out of memory.
3. **Stateful services.** Prometheus, Grafana and Loki are pinned to the manager (`node.role == manager`) so their named volumes always re-attach to the same host.

If we need more capacity, we just update the `worker_count` variable in Terraform to scale horizontally or to spin up more droplets, then use `docker service scale minitwit_stack_minitwit=N` to add more copies of the app.

3. Reflection Perspective

3.1 Database dataloss during first migration from SQLite to MySQL

During the migration from SQLite to the DigitalOcean MySQL database, our biggest issue was unexpected data loss: the new database was mostly empty, while the simulator still expected the historical old users and follow relationships to exist. This caused steadily increasing errors in the tweet, follow, and unfollow endpoints, even though registration still worked. After checking Docker, API behavior, and Grafana metrics, we found that the real problem was not the code or infrastructure performance, but missing production data.

We solved the issue by writing a custom SQLAlchemy migration script that connected to both the old SQLite database and the new MySQL database, then restored the critical User and Follower tables. Around 98% of the data was recovered, although some parts were still missing. This still caused a slow error increase but we managed to scale down the issue drastically. The main lessons learned are that database migrations must include proper backups, validation checks, and rollback plans. At the end we learned that nothing is as important as data already in production.

3.2 Slow public timeline query and user-related pages

Our second major issue was the extremely slow public timeline and user-related pages, where loading times sometimes reached 30 seconds to 1 minute. At first, we thought the problem was caused by high memory usage from the new Docker Swarm setup, especially because we started the swarm on 1 droplet which we scaled up horizontally. After investigation into the sql and database setup we found that queries took up all memory and required crazy amount of resources that would be needed for other services. The first fix was to introduce database indexes, which improved the public timeline queries, but it did not fully solve login and other user-related operations. The deeper problem was the way our queries were formatted, especially queries using OR, which forced inefficient database lookups and created unnecessary memory pressure. Also have to note here the minimalist droplet database hardware setup is not designed for 4 million messages and 100k user setup.

We solved the issue by changing the query structure instead of only relying on indexes. In particular, we removed the expensive OR statement and first created the relevant list of users before running the final query. This made the database access more predictable and reduced the load on the droplets. The main lesson learned is that performance problems are not always fixed by

scaling infrastructure or adding indexes; query design and data access patterns are just as important. We learned that for future we need to monitor slow endpoints, inspect actual database queries, and treat performance refactoring as part of normal system maintenance rather than only reacting when pages become unusable.

3.3 Loki timeout error under memory pressure

Another major issue we faced was with logging in the observability stack, specifically Loki under memory pressure. Promtail on the manager node started failing with context deadline exceeded when trying to push logs to the Loki API. The root cause was that the Loki ingester kept log chunks in memory for too long. Both `chunk_idle_period` and `max_chunk_age` were set to 1 hour. With Loki limited to only 280 MB of memory, this caused heavy garbage-collection pauses, which made Promtail time out and temporarily broke reliable log collection.

We solved it by tuning Loki's memory behavior. Both chunks were reduced from 1 hour to 10 minutes in `monitoring/loki-config.yaml`, and Loki's memory limit was increased from 280 MB to 420 MB in `docker-compose.yml`. The main lesson learned is that monitoring and logging services also need proper resource planning.

3.4 “DevOps” style of work

The DevOps style of our work was different from previous development projects because we did not only focus on writing application features. We also had to think about deployment, infrastructure, monitoring, logging, performance, and recovery as part of the same development process. Instead of manually running the app and checking if it worked locally, we used Docker, Docker Swarm, GitHub Actions, Prometheus, Grafana, Loki, and DigitalOcean to build a production-like system. The hardest part was managing all these different components which we haven't used before to make the separate work together seamlessly. The biggest issue definitely we had with the database as seen in the 3.1 and 3.2 paragraphs. Since we thought for so long that the application and the droplets were causing issues. This highlights how important a correct and well put together monitoring/logging system is the hearth of an application.

4. Use of Generative AI

In accordance with ITU's guidelines on the use of generative AI for assessed work, we disclose the following.

Tools used. Claude (Anthropic) and Gemini - primarily Claude. Both were used as pair-programming and writing assistants, not as autonomous agents: every diff went through human review and the CI quality gate before merging.

Where they helped.

- **YAML and shell.** First drafts of `docker-compose.yml` placement constraints, Prometheus DNS-SD scrape configs, Grafana dashboard JSON skeletons and the `deploy.sh` hashing logic were AI-assisted, then audited line-by-line.
- **Documentation and report.** The structure of `SECURITY.md`, of `docs/infrastructure-as-code.md` and of this report was drafted with an LLM and then revised against the actual source files. Wording polish and consistency checks were AI-assisted.
- **Bug diagnosis.** Throught the process we faced some database errors and docker firewall issues that were diagnosed faster by walking an LLM through the symptoms and asking it to enumerate hypotheses and solutions;

Co-authorship. LLM tooling is registered as co-author LLM `<none>` in `.mailmap` per the assignment instructions.

Word count: ~2,300 words (excluding tables, code blocks, headings and front-matter).